

CHINMAY DATTANAND KUCHINAD

+1(213) 378-5916 ckuchina@usc.edu [in](#) Chinmay Dattanand Kuchinad [G](#) chinmaydk99

EDUCATION

University of Southern California, Los Angeles

Master of Science - Specialisation in Machine Learning and Data Science

Aug 2023 – May 2025

GPA: 3.61/4.00

PES University, Bengaluru

Bachelors of Technology in Electronics and Communication Engineering

Aug 2016 – Jul 2020

GPA: 8.84/10.00

EXPERIENCE

USC Information Sciences Institute

Researcher – Edge AI & Distributed Machine Learning

Feb 2025 – May 2025

Los Angeles, California

- Developing **pipeline parallelism strategies for the PipeEdge framework**, creating dynamic partitioning algorithms that **intelligently allocate model layers** based on compute capabilities and network bandwidth across heterogeneous edge devices, **reducing end-to-end latency by 43%** compared to conventional data-parallel approaches
- Implementing both **post-training quantization and magnitude-based pruning methods** with iterative fine-tuning, achieving 60% sparsity in Vision Transformer(Base16) models through **structured channel pruning while preserving critical features for inference tasks**

Samsung Electronics America

Generative AI Application Intern

Jun 2024 – Aug 2024

Ridgefield Park, New Jersey

- Developed a **multi-agent email campaign generation system** using open-source LLMs to orchestrate specialized agents for email marketing campaigns via LangGraph, ensuring consistent brand voice through advanced prompt engineering
- Engineered a **ChromaDB-powered RAG system** for precise product information retrieval, ensuring grounded responses and reducing hallucinations
- Fine-tuned open source LLMs (Llama-3.1, Gemma) on historical marketing campaigns data using PEFT and QLoRA techniques

Bosch Global Software Technologies Pvt Ltd

Software Engineer - Machine Learning

Jan 2020 – Nov 2022

Bengaluru, India

- Built production ML pipelines (PyTorch) for Honda Japan, analyzing operational data for engine performance optimization; **reduced RAM usage prediction error by 32%**, enhancing predictive maintenance
- Developed an **automated FOTA testing framework**, leveraging Python and SQL to perform in-depth analysis of device telemetry data. Built ML models that **increased flash success rates from 73% to 91%** by identifying key factors influencing success and optimizing for variable network conditions, significantly improving system uptime and reducing update failures
- Formulated and managed **data pipelines orchestrated for the ingestion of ECU(Engine Control Unit) data**, overseeing insights such as EEPROM usage, fault prediction, and memory usage

Projects

AI-Powered Multi-Agent System for Sales Automation

Orchestrating LLMs using OpenAI agents sdk for End-to-End Sales Email Automation with Guardrails

April 2025 – May 2025

- Architected and built an asynchronous, **multi-agent AI system in Python automating end-to-end sales email generation**, integrating LLMs (Gemini, GPT-4o-mini) and SendGrid via OpenAI agentic SDK framework **featuring Pydantic validation and guardrails** (professionalism, brand compliance), designed to scale personalized SDR outreach and enhance communication consistency

LLM-Based Clinical Diagnosis Prediction

Fine-tuning Llama-3.2 on MIMIC-IV dataset for Next-Visit Disease Forecasting

Aug 2024 – Dec 2024

- Developed an LLM-powered **sequential disease forecasting model by fine-tuning Llama-3.2** with QLoRA and 4-bit quantization on the MIMIC-IV EHR dataset, achieving **85.43% ROC-AUC and 42.32% PR-AUC** in predicting next-visit diagnoses and demonstrating efficient, high-fidelity clinical predictive modeling

TECHNICAL SKILLS

Programming Scripting: Python, C++, R, SQL, JavaScript, TypeScript, CUDA

Machine Learning and Deep Learning: PyTorch, TensorFlow, Hugging Face, Scikit-learn, XGBoost, LightGBM, MLflow, Grafana, Prometheus

Data Visualization & Engineering: Pandas, NumPy, SciPy, Matplotlib, Seaborn, Plotly, Kafka, Spark, Airflow, Snowflake

Cloud & Backend: AWS (SageMaker, S3, EKS, Lambda), Docker, Kubernetes, Node.js, Express.js, FastAPI, Flask, Github Actions

GenAI Tools & Frameworks: Cursor, Bolt, MCP, Langchain, Langgraph, CrewAI, OpenAI Agents SDK, Google ADK